

Ilyass Ouardi

M.Sc. Student in Computer Science
(+39) 320 175 5208 · ilyassouardi@gmail.com
[ouardi-ilyass.github.io](https://github.com/ouardi-ilyass) · [GitHub](#) · [LinkedIn](#)

Research Interests

I am interested in **AI safety** and **mechanistic interpretability**. To me, interpretability is not an academic exercise for its own sake, but a pragmatic tool to make model safety measurable and verifiable on the path toward safe AGI. Right now I am particularly focused on the **geometry of representations**: how abstract concepts — like refusal and truth — are encoded within a model’s activation space. Two questions follow from this. The first is one of *model biology*: whether such concepts conform strictly to single linear directions, or rely instead on higher-dimensional, possibly non-linear manifolds. The second is about *applied interpretability*: whether representation-level edits — steering, ablation, concept erasure — are genuinely surgical when evaluated against practical baselines, with off-target side-effects measured as rigorously as the primary behavioral change. The two are connected, since an intervention built on the wrong picture of the geometry is precisely one whose side-effects go unmeasured.

Education

M.Sc. in Computer Science, University of Milan, Italy Sep 2024 – Present

Thesis: *The Geometry of Absence, Opposite, and the Activation Manifold in Representation-Level Interventions on LLMs*
Advisors: Prof. Alfio Ferrara & Elisabetta Rocchetti (University of Milan)

Focus: Investigating representation-level interventions (linear concept erasure, steering) and studying how models encode the absence of a concept versus its opposite on the activation manifold.

Expected defense: Summer 2027 (July 2027) · Weighted average: **29.22/30** (3.90/4.00)

Selected coursework: Natural Language Processing (30 *cum laude*), Distributed and Pervasive Systems (30 *cum laude*), Security of Service-Oriented Architectures (30 *cum laude*), Advanced Biometrics (30), Audio Pattern Recognition (30), Enterprise Security Management (30), Privacy and Data Protection (29)

B.Sc. in Computer Science, University of Milan, Italy

Sep 2020 – Oct 2024

Thesis: *Design and development of a donation-tracking system for RECUP APS*

Advisor: Prof. Carlo Maria Bellettini

Final grade: **110/110 cum laude**; weighted average **28.55/30** (3.81/4.00)

Mathematical and theoretical foundations: Discrete Mathematics (30 *cum laude*), Mathematical Logic (30 *cum laude*), Automata and Formal Languages (30 *cum laude*), Programming Languages (30 *cum laude*), Artificial Intelligence I (30 *cum laude*), Calculus (30), Statistics and Data Analysis (30), Algorithms and Data Structures (30)

Diploma, Computer Science (Perito Informatico), “G. e M. Montani” Institute, Fermo

Sep 2015 – Jul 2020

Final grade: **100/100 cum laude**

Research Experience

Independent Research — Representation Geometry of Truth, ISLab, University of Milan Dec 2025 – Jul 2026

Sole author. Natural Language Processing course.

- Tested whether propositional truth is mediated by one linear direction or by a k -dimensional concept cone. Built **TDO** and **TCO**: gradient-optimised directions and cones under causal necessity, sufficiency and KL-retention losses.
- **Result**: a single direction is insufficient for the larger models, and a low-dimensional cone ($k=2$) is the better model of the concept. Its basis is near-orthogonal to the Difference-in-Means probe yet individually causally effective.
- **Manuscript**: *The Multidimensional Geometry of Truth in LLMs: A Concept-Cone Extension of the Linear Representation Hypothesis* (2026). [\[Draft\]](#) [\[Code\]](#)

Research Project — What a Heart-Sound Classifier Actually Learns, University of Milan Apr 2026 – Jul 2026
Sole author. Audio Pattern Recognition course.

- Compared three feature representations against three classifier families on PhysioNet/CinC 2016 under one recording-level protocol, pairing every performance claim with a diagnostic capable of refuting it.
- **Result:** the headline MAcc of 0.866 is largely a recording-site artefact — site recoverable from the same features at 0.647 against 0.167 chance, diagnosis collapsing from 0.840 to 0.457 under leave-one-site-out. Attribution leans systolic, but no more so on correct than on incorrect predictions.
- **Manuscript:** *What Does a Heart Sound Classifier Actually Learn?* (2026). [Draft] [Code]

Research Project — Empirical Test of the Legends Hypothesis, SESAR Lab, University of Milan Feb 2026 – Jun 2026

With F. Bylyshi & H. El-Khazri (equal contribution)

- Tested whether a model learns a regulation better from stories showing it being followed than from the regulation itself. Built a pipeline producing training corpora matched on everything but content, and **GenderEqGLUE**, a five-task benchmark evaluated on documents held out from training.
- **Result:** the two regimes teach complementary competences — training on the rule wins overall, training on the stories wins on the tasks that require applying it.
- **Manuscript:** *The Legend Challenge: Embedding Ethical Compliance into LLMs through Champion Narratives* (2026). [Draft] [Code]

Teaching Experience

Teaching Assistant, *Accertamento delle Competenze Informatiche* (3 CFU) Jan 2026 – Present
University of Milan, Scientific lead: Prof. Stefano Montanelli; delivered by the CTU

- Prepared and delivered lectures to classes of students, authoring the accompanying teaching material.
- Led lab sessions of a blended-learning course shared across numerous undergraduate degree programmes, and invigilated examinations.

Engineering Experience

Lead Developer / Consultant, RECUP APS – University of Milan Jan 2024 – Dec 2025

- Designed and built a donation-tracking web platform (React + TypeScript, Supabase DB/Auth/Storage, Netlify Functions); began as B.Sc. thesis, continued as a consultancy with me as lead developer.
- **RecupGPT** (2025): retrieval-augmented natural-language querying over the production database (LangChain, LangGraph, NL2SQL with query rewriting and routing).

Full-Stack Developer Consultant, Value Process Srl, Milan Jun 2022 – Sep 2022

- Microservices architecture in Spring Boot and React; led migration of a legacy system to a serverless architecture.

Fellowships, Awards, and Honours

Full academic + merit scholarships (tuition and housing), University of Milan Sep 2020 – Sep 2026

Lead the Future — STEM mentorship community, Mentee Sep 2025 – Present

BOOST '24 — selected for the Bologna Orthogonal Summer Term (outstanding B.Sc./M.Sc. CS students) Aug 2024

Futuro Annunciato '24 — selected for the event, for outstanding students Jul 2024

Nomination, “Alfieri del Lavoro” Award, Presidency of the Italian Republic Jul 2020

Participant, National Olympiads in Informatics & Mathematics 2017 – 2020

Technical Skills

Mechanistic interpretability: denoising activation patching, directional ablation and activation addition, Difference-in-Means probes, gradient-based direction/cone optimisation, concept erasure (INLP, LEACE), activation steering

Post-hoc XAI: SHAP (tree and kernel estimators), Grad-CAM with saliency sanity checks, structured occlusion

Experimental method: causal mediation protocols, permutation and analytic nulls, grouped / leave-one-group-out cross-validation, paired significance testing with multiple-comparison correction, bootstrap over the correct unit of independence

ML & LLM tooling: PyTorch, nnsight, scikit-learn, LangChain, LangGraph, RAG, NL2SQL, supervised fine-tuning

Programming — *proficient:* Python, Java, TypeScript; *familiar:* Go, C, Scala, OCaml, Erlang

Web, data & cloud: React, SQL / PL-SQL, Supabase (PostgreSQL + RBAC), Netlify Functions, CI/CD

Languages: Italian (native), Arabic (native), English (fluent, working language)

Last updated: October 5, 2026